

# CSCE 421: Machine Learning

## Lecture 11: Random Forest and Boosting

Texas A&M University

CSCE 421

Bobak Mortazavi

Ryan King

Zhale Nowroozilarki

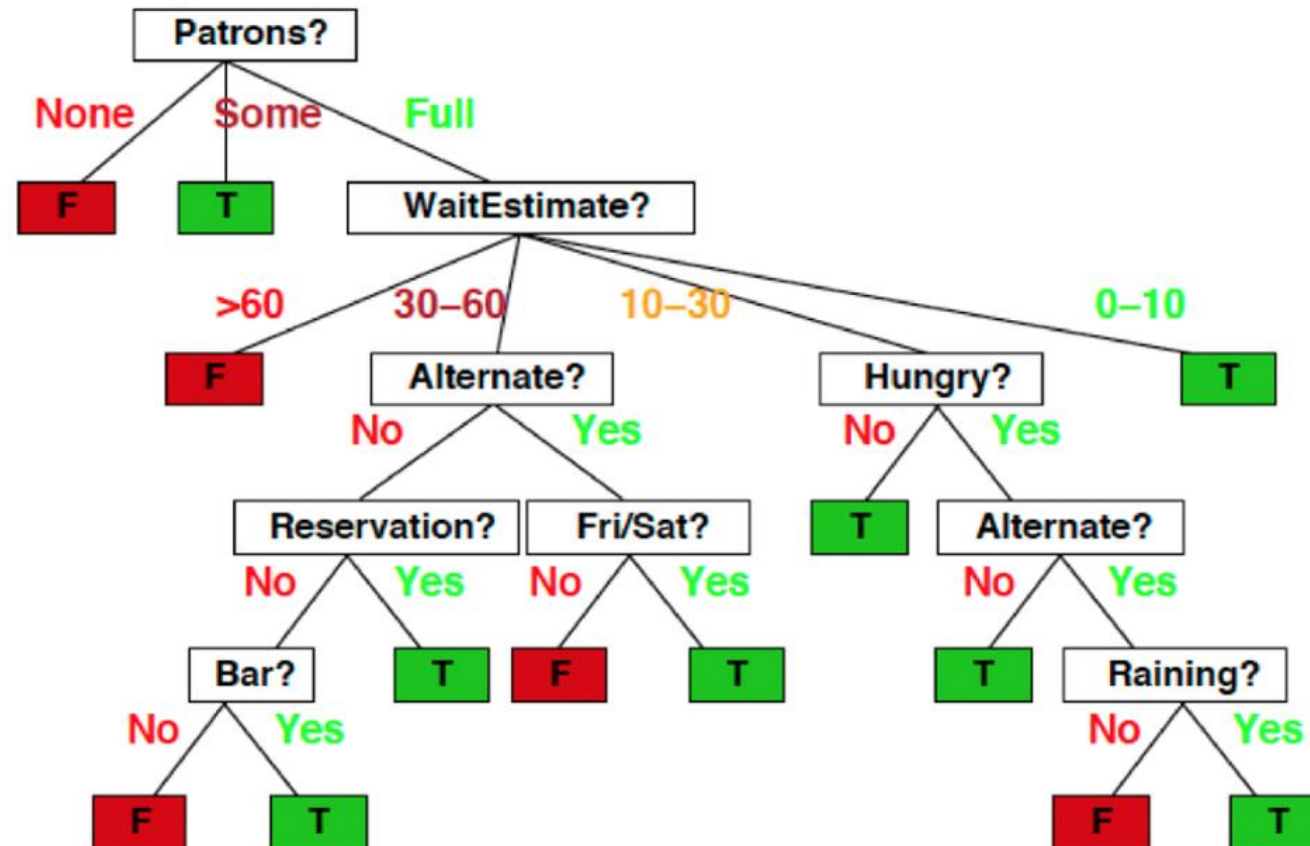
# Goals

---

- Review of decision trees
- Overcoming the limitations to decision trees
- Introduce random forests
- Introduce gradient descent boosting

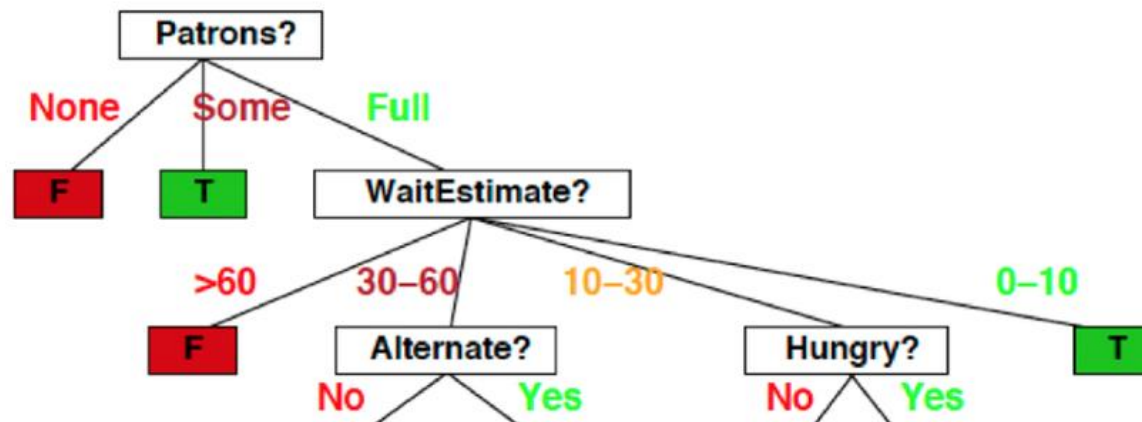
# Final decision tree

Greedily we build the tree and looks like this



# Pruning classification trees

We should **prune** some of the leaves of the tree to get a smaller depth



- If we stop here, not all training samples are classified correctly
- **How do we classify a new instance?**
  - We label the leaves of this smaller tree with the label of the majority of training samples

# Decision trees vs. other models

- Advantages:
  - Models are transparent: easily interpretable!
  - Data can contain combination of feature types: Qualitative predictors without dummy variables
  - Decision trees more closely mirror human decision making
  - Graphical representation
- Disadvantages:
  - Usually not same level of predictive accuracy
  - Not robust (small change in the data can change the tree a lot)

# Bagging: Bootstrapped Aggregating

---

- Accuracy of Decision Trees suffer from high variance
- For example if we split data in half, tree for both halves could be very different

# Tree 1: Restaurants

---

## Tree 2: Restaurants

---



## Bagging: Why does it work?

- Given  $n$  independent observations  $Z_1, \dots, Z_n$  each with variance  $\sigma^2$
- Variance of the mean  $\bar{Z} = \frac{\sigma^2}{n}$
- That's lower variance!
- So, how do we take advantage of this?

# Bagging Trees

- Take  $B$  different training sets
- Train  $f^1$  on training set 1
- Train  $f^2$  on training set 2
- ...

# Bagging Trees

- Take B different training sets
  - Train  $f^1$  on training set 1
  - Train  $f^2$  on training set 2
  - ...
- 
- Can average the result over B trees, as a single, low-variance model

$$\hat{f}_{avg}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$$

# Bagging Trees

- Take B different training sets
  - Train  $f^1$  on training set 1
  - Train  $f^2$  on training set 2
  - ...
- 
- Can average the result over B trees, as a single, low-variance model

$$\hat{f}_{avg}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x)$$

- But where do we come up with B training sets??

# Bootstrapping

- Take B different bootstraps of our one datasets
- $\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^{B^*} \hat{f}^b(x)$
- Turns out, you can grow these trees without pruning them
- For Regression – average the values from each tree
- For Classification – take the majority vote across trees
- Test error can be plotted as a function of B

# Bootstrapping

- Take B different bootstraps of our one datasets
  - $\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^{B^*} \hat{f}^b(x)$
  - Turns out, you can grow these trees without pruning them
  - For Regression – average the values from each tree
  - For Classification – take the majority vote across trees
  - Test error can be plotted as a function of B
- 
- Great Fact: B turns out not to be a critical parameter, so large B does not mean we overfit!

# Measuring Error

- If we repeatedly fit bootstrapped subsets (say  $2/3$ ) of our data
- Each time we will be left with a subset (say  $1/3$ ) that we can call out of bag
- We can then estimate the error for this as we train – we call this the Out of Bag Estimation

# New Example: Heart Dataset

<https://archive.ics.uci.edu/ml/datasets/heart+Disease>

```
> summary(data)
```

X	Age	Sex	ChestPain	RestBP	chol	Fbs
Min. : 1.0	Min. :29.00	Min. :0.0000	asymptomatic:144	Min. : 94.0	Min. :126.0	Min. :0.0000
1st Qu.: 76.5	1st Qu.:48.00	1st Qu.:0.0000	nonanginal : 86	1st Qu.:120.0	1st Qu.:211.0	1st Qu.:0.0000
Median :152.0	Median :56.00	Median :1.0000	nontypical : 50	Median :130.0	Median :241.0	Median :0.0000
Mean :152.0	Mean :54.44	Mean :0.6799	typical : 23	Mean :131.7	Mean :246.7	Mean :0.1485
3rd Qu.:227.5	3rd Qu.:61.00	3rd Qu.:1.0000		3rd Qu.:140.0	3rd Qu.:275.0	3rd Qu.:0.0000
Max. :303.0	Max. :77.00	Max. :1.0000		Max. :200.0	Max. :564.0	Max. :1.0000


RestECG	MaxHR	EXAng	oldpeak	slope	Ca	Thal
Min. :0.0000	Min. : 71.0	Min. :0.0000	Min. :0.00	Min. :1.000	Min. :0.0000	fixed : 18
1st Qu.:0.0000	1st Qu.:133.5	1st Qu.:0.0000	1st Qu.:0.00	1st Qu.:1.000	1st Qu.:0.0000	normal :166
Median :1.0000	Median :153.0	Median :0.0000	Median :0.80	Median :2.000	Median :0.0000	reversible:117
Mean :0.9901	Mean :149.6	Mean :0.3267	Mean :1.04	Mean :1.601	Mean :0.6722	NA's : 2
3rd Qu.:2.0000	3rd Qu.:166.0	3rd Qu.:1.0000	3rd Qu.:1.60	3rd Qu.:2.000	3rd Qu.:1.0000	
Max. :2.0000	Max. :202.0	Max. :1.0000	Max. :6.20	Max. :3.000	Max. :3.0000	

AHD
No :164
Yes:139



# Comparison of Algorithms



# Variable Importance: Multiple Trees


---

## Variable Importance: Lost

- Interpreting Bagging is difficult
- No longer possible to decide variable order from a single tree
- With regression trees – must understand summary reduction in RSS at each split
- With classification trees – overall summary in reduction in Gini Index or Entropy at each split
- Can create a new term: Relative Importance

$$v_j = \frac{1}{M} \sum_{m=1}^M I(j \in T_m)$$

# Heart Dataset Variable Importance



# Bagging Limitations: Strong Predictor

---


# Bagging Limitations

- What if you have a strong predictor and then a bunch of moderate predictors?
- Each time, the first variable is that strong predictor!
- So, are these trees really any different? In other words, is variance really reduced?
- What if at each split of each tree we only consider a subset  $m$  of predictors  $p$ ?
- In other words: What if we randomly eliminate the strong predictor when making some of the trees?

# Random Forests


- Set  $m \approx \sqrt{p}$
- Each time you want to build a tree:  $\frac{p-m}{p}$  predictors aren't considered
- This gives other moderate predictors a chance to be important!
- The average tree becomes less variable and thus, more reliable!

# Random Forest: Heart Disease





## Random Forest: Adjusting m



# What we have learnt so far!

- Decision Trees
  - Hierarchical structure to perform modeling
  - Tree structure determined by splitting criterion
  - Pruning: prevents overfitting by limiting depth of the tree
  - Main advantage: interpretable!
- Random Forests
  - Ensemble of trees
  - Through bagging and randomization results in reduced variance
  - Great performance in practice
  - Reduced interpretability

# What we have learnt so far!

- Decision Trees
  - Hierarchical structure to perform modeling
  - Tree structure determined by splitting criterion
  - Pruning: prevents overfitting by limiting depth of the tree
  - Main advantage: interpretable!
- Random Forests
  - Ensemble of trees
  - Through bagging and randomization results in reduced variance
  - Great performance in practice
  - Reduced interpretability

CAN WE DO EVEN BETTER?

# Boosting

- We can potentially improve decision tree performance even further
  - Develop a method here that actually works on any classifier
  - Decision Tree: Build a tree based upon variable importance in training data
  - Bagging: Build each tree randomly – reduces variance, improves overfitting issues
  - Random Forest: build each tree randomly, with random variations in predictors, improves performance/accuracy
- 
- So: What if we build trees in a sequential, ordered fashion?

# Generalized Boosting

- Can we build a high-capacity model “one unit at a time”?

$$model(x, w) = w_0 + f_1(x) * w_1 + f_2(x) * w_2 + ... + f_M(x) * w_M$$



# Generalized Boosting

- Can we build a high-capacity model “one unit at a time”?

$$model(x, w) = w_0 + f_1(x) * w_1 + f_2(x) * w_2 + ... + f_M(x) * w_M$$

- The basic principle of boosting is to progressively build a high capacity model one unit at a time

# Boosting Capacity



# Sequential Building

- Set of M nonlinear features or units from a single family of universal approximators

$$F = \{f_1(x), f_2(x), \dots, f_M(x)\}$$

- Add units sequentially (one at a time) to build a set of M models that increase in complexity with respect to the trained data, ending with a generic nonlinear model composed of M units



# Boosted Model

$$model(x, w) = w_0 + f_{s_1}(x) * w_1 + f_{s_2}(x) * w_2 + ... + f_{s_M}(x) * w_M$$

- Re-indexed the individual units to  $f_{s_m}$  to denote the unit from the entire collection  $F$  added in the  $m$ th round
- The linear combination of weights  $w$  are collectively represented sometimes as  $\theta$

# Boosting Procedure

- Process of boosting is performed for a total of  $M$  rounds
- At each round, we determine which unit, when added to the running model, best lowers its training error
- We then measure the corresponding validation error
- For the sake of simplicity – let's stick with regression for now
- Classification remains exactly the same!

## Initial Round (Round 0)

- Start with model:

$$\text{model}_0(x, \theta) = w_0$$

- Whose weight set  $\theta_0 = \{w_0\}$ , which contains a single bias weight which minimizes least squares (notation from Watt, Borhani, and Kastagelos – P is people):

$$\frac{1}{p} \sum_{p=1}^P (\text{model}_0(x_p, \theta_0) - y_p)^2 = \frac{1}{p} \sum_{p=1}^P (w_0 - y_p)^2$$

- This optimal  $w_0$  remains fixed forever forward

## Round 1 of Boosting

- Having tuned the only parameter, we now boost its complexity by adding weighted unit  $f_{s_1}(x) w_1$  :

$$model_1(x, \theta_1) = model_0(x, \theta_0) + f_{s_1}(x) w_1$$

- To determine which unit in our set  $F$  best lowers the training error, we pick the  $f_s \in F$  that minimizes the cost:

$$\frac{1}{p} \sum_{p=1}^P (model_0(x_p, \theta_0) + f_{s_1}(x_p) w_1 - y_p)^2 =$$
$$\frac{1}{p} \sum_{p=1}^P (w_0 + f_{s_1}(x_p) w_1 - y_p)^2$$

## Round $m > 1$ of Boosting

$$model_{m-1}(x, \theta_{m-1}) = w_0 + f_{s_1}(x) * w_1 + f_{s_2}(x) * w_2 + \dots + f_{s_{m-1}}(x) * w_{m-1}$$

- We then seek out the best next unit to add

$$model_m(x, \theta_m) = model_{m-1}(x, \theta_{m-1}) + f_{s_m}(x) w_m$$

- By minimizing


$$\begin{aligned} \frac{1}{p} \sum_{p=1}^P (model_{m-1}(x_p, \theta_{m-1}) + f_{s_m}(x_p) w_m - y_p)^2 = \\ \frac{1}{p} \sum_{p=1}^P (w_0 + f_{s_1}(x_p) w_1 + \dots + f_{s_m}(x_p) w_m - y_p)^2 \end{aligned}$$

# Boosting Loss

$$\frac{1}{p} \sum_{p=1}^P (w_0 + f_{s_1}(x_p) w_1 + \dots + f_{s_m}(x_p) w_m - y_p)^2$$

- If we use a fixed-shape approximator, say a decision stump or decision tree, this entails solving  $M$  (or  $M - m + 1$ , if we decide to check only those units not used in previous rounds) optimization problems
- If we use neural networks, each unit takes the same form, we only need to solve one such optimization problem (for future reference)
- Once boosting is complete we select from our set of models the one that provides the lowest validation error
- Alternatively, can halt if the validation error increases (overfits) – \*early stopping\*
- Be careful, validation error can oscillate

# Visualization using trees



# Boosting for Regression Trees: Algorithm

1. Set  $\hat{f}(x) = 0$  and error  $r_i = y_i$



# Boosting for Regression Trees: Algorithm

1. Set  $\hat{f}(x) = 0$  and error  $r_i = y_i$
2. For  $b = 1, 2, \dots, B$  repeat:
  - a. Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d + 1$  terminal nodes) to the training data  $(X, r)$
  - b. Update  $\hat{f}$  by adding in a shrunk version of the new tree

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$

- c. Update the residuals

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$$

# Boosting for Regression Trees: Algorithm

1. Set  $\hat{f}(x) = 0$  and error  $r_i = y_i$
2. For  $b = 1, 2, \dots, B$  repeat:
  - a. Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d + 1$  terminal nodes) to the training data  $(X, r)$
  - b. Update  $\hat{f}$  by adding in a shrunk version of the new tree

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$

- c. Update the residuals

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$$


3. Output the boosted model

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

# Boosted Decision Trees

- Learn slowly from shallow trees
- Given a current model – calculate residuals
- Build next tree to improve on the remaining residuals
- Slowly improve where the model does not currently perform well
- Boosted classification becomes a bit trickier in how it updates (next time)

# Boosting vs. Random Forest



# What we have learnt so far!

- Decision Trees
  - Hierarchical structure to perform modeling
  - Tree structure determined by splitting criterion
  - Pruning: prevents overfitting by limiting depth of the tree
  - Main advantage: interpretable!
- Random Forests
  - Ensemble of trees
  - Through bagging and randomization results in reduced variance
  - Great performance in practice
  - Reduced interpretability
- Boosted Trees
  - Ensemble of trees (boosts a bunch of weak classifiers into a strong classifier)
  - Through sequential building, improves model performance
  - If  $B$  is too large, this model Does overfit
  - $\lambda$  is small, but greater than 0
  - Depth  $d$  of trees is often small ( $d = 1$ , decision stumps, are very interpretable)

## Next Time

---

- More Boosting
- Gradient Descent Boosting for Classification
- Loss Functions
- Python Coding
- QUIZ 2!